

# Scatter 层[TODO]

- mode (要求 TensorRT>=8.2)
  - Scatter ELEMENT 模式
  - Scatter ND 模式

## mode (要求 TensorRT>=8.2)

### ELEMENT 模式

```
import numpy as np
from cuda import cudart
import tensorrt as trt

np.random.seed(97)
nIn, cIn, hIn, wIn = 1, 3, 4, 5 # 输入张量 NCHW
data0 = np.arange(nIn * cIn * hIn * wIn, dtype=np.float32).reshape(nIn, cIn, hIn, wIn) # 输入数据
data1 = np.tile(np.arange(hIn), [nIn, cIn, 1, wIn]).astype(np.int32).reshape(nIn, cIn, hIn, wIn)
data2 = -data0

np.set_printoptions(precision=8, linewidth=200, suppress=True)
cudart.cudaDeviceSynchronize()

# 仅适用于 axis = 2
def scatterCPU(data0, data1, data2):
    nIn, cIn, hIn, wIn = data0.shape
    output = data0
    for n in range(nIn):
        for c in range(cIn):
            for h in range(hIn):
                for w in range(wIn):
                    output[n, c, data1[n, c, h, w], w] = data2[n, c, h, w]
                    #print("<%(n,c,h,w,n,c,data1[n,c,h,w],w,data2[n,c,h,w]))
                    #print(output)
    return output

logger = trt.Logger(trt.Logger.ERROR)
builder = trt.Builder(logger)
network = builder.create_network(1 << int(trt.NetworkDefinitionCreationFlag.EXPLICIT_BATCH))
config = builder.create_builder_config()
config.max_workspace_size = 1 << 30
inputT0 = network.add_input('inputT0', trt.DataType.FLOAT, (nIn, cIn, hIn, wIn))
inputT1 = network.add_input('inputT1', trt.DataType.INT32, (nIn, cIn, hIn, wIn))
inputT2 = network.add_input('inputT2', trt.DataType.FLOAT, (nIn, cIn, hIn, wIn))
#-----# 替换部分
scatterLayer = network.add_scatter(inputT0, inputT1, inputT2, trt.ScatterMode.ELEMENT)
scatterLayer.axis = 2
#-----# 替换部分
network.mark_output(scatterLayer.get_output(0))
engineString = builder.build_serialized_network(network, config)
engine = trt.Runtime(logger).deserialize_cuda_engine(engineString)
context = engine.create_execution_context()
_, stream = cudart.cudaStreamCreate()
```

```

inputH0 = np.ascontiguousarray(data0.reshape(-1))
inputH1 = np.ascontiguousarray(data1.reshape(-1))
inputH2 = np.ascontiguousarray(data2.reshape(-1))
outputH0 = np.empty(context.get_binding_shape(3), dtype=trt.nptype(engine.get_binding_dtype(3)))
_, inputD0 = cudart.cudaMallocAsync(inputH0.nbytes, stream)
_, inputD1 = cudart.cudaMallocAsync(inputH1.nbytes, stream)
_, inputD2 = cudart.cudaMallocAsync(inputH2.nbytes, stream)
_, outputD0 = cudart.cudaMallocAsync(outputH0.nbytes, stream)

cudart.cudaMemcpyAsync(inputD0, inputH0.ctypes.data, inputH0.nbytes,
cudart.cudaMemcpyKind.cudaMemcpyHostToDevice, stream)
cudart.cudaMemcpyAsync(inputD1, inputH1.ctypes.data, inputH1.nbytes,
cudart.cudaMemcpyKind.cudaMemcpyHostToDevice, stream)
cudart.cudaMemcpyAsync(inputD2, inputH2.ctypes.data, inputH2.nbytes,
cudart.cudaMemcpyKind.cudaMemcpyHostToDevice, stream)
context.execute_async_v2([int(inputD0), int(inputD1), int(inputD2), int(outputD0)], stream)
cudart.cudaMemcpyAsync(outputH0.ctypes.data, outputD0, outputH0.nbytes,
cudart.cudaMemcpyKind.cudaMemcpyDeviceToHost, stream)
cudart.cudaStreamSynchronize(stream)

print("inputH0 :", data0.shape)
print(data0)
print("inputH1 :", data1.shape)
print(data1)
print("inputH2 :", data2.shape)
print(data2)
print("outputH0:", outputH0.shape)
print(outputH0)

cudart.cudaStreamDestroy(stream)
cudart.cudaFree(inputD0)
cudart.cudaFree(outputD0)

resCPU = scatterCPU(data0,data1,data2)
print("diff:\n",outputH0 - resCPU)

```

- 输入张量 0 形状 (1, 3, 4, 5), 输入张量 2 形状与输入张量 0 相同, 值为其相反数

$$\left[ \left[ \begin{bmatrix} 0. & 1. & 2. & 3. & 4. \\ 5. & 6. & 7. & 8. & 9. \\ 10. & 11. & 12. & 13. & 14. \\ 15. & 16. & 17. & 18. & 19. \end{bmatrix} \begin{bmatrix} 20. & 21. & 22. & 23. & 24. \\ 25. & 26. & 27. & 28. & 29. \\ 30. & 31. & 32. & 33. & 34. \\ 35. & 36. & 37. & 38. & 39. \end{bmatrix} \begin{bmatrix} 40. & 41. & 42. & 43. & 44. \\ 45. & 46. & 47. & 48. & 49. \\ 50. & 51. & 52. & 53. & 54. \\ 55. & 56. & 57. & 58. & 59. \end{bmatrix} \right] \right]$$

- 输入张量 1 形状 (1, 3, 4, 5)

$$\left[ \left[ \begin{bmatrix} 0 & 1 & 2 & 3 & 0 \\ 1 & 2 & 3 & 0 & 1 \\ 2 & 3 & 0 & 1 & 2 \\ 3 & 0 & 1 & 2 & 3 \end{bmatrix} \begin{bmatrix} 0 & 1 & 2 & 3 & 0 \\ 1 & 2 & 3 & 0 & 1 \\ 2 & 3 & 0 & 1 & 2 \\ 3 & 0 & 1 & 2 & 3 \end{bmatrix} \begin{bmatrix} 0 & 1 & 2 & 3 & 0 \\ 1 & 2 & 3 & 0 & 1 \\ 2 & 3 & 0 & 1 & 2 \\ 3 & 0 & 1 & 2 & 3 \end{bmatrix} \right] \right]$$

- 输出张量 0 形状 (1, 3, 4, 5)

$$\left[ \left[ \begin{bmatrix} -0. & -16. & -12. & -8. & -4. \\ -5. & -1. & -17. & -13. & -9. \\ -10. & -6. & -2. & -18. & -14. \\ -15. & -11. & -7. & -3. & -19. \end{bmatrix} \begin{bmatrix} -20. & -36. & -32. & -28. & -24. \\ -25. & -21. & -37. & -33. & -29. \\ -30. & -26. & -22. & -38. & -34. \\ -35. & -31. & -27. & -23. & -39. \end{bmatrix} \begin{bmatrix} -40. & -56. & -52. & -48. & -44. \\ -45. & -41. & -57. & -53. & -49. \\ -50. & -46. & -42. & -58. & -54. \\ -55. & -51. & -47. & -43. & -59. \end{bmatrix} \right] \right]$$

- 含义: 参考 [Onnx ScatterElements 算子](#) 和 [TensorRT C++ API 说明](#)

- 数据张量 `data`、索引张量 `index`、更新张量 `update`、输出张量 `output` 形状相同 ( $dim = r$ )，均为  $[d_0, d_1, \dots, d_{r-1}]$ ，指定  $axis = p$  ( $0 \leq p < r$ )，则
- 命循环变量  $i_j$  满足  $0 \leq j < r, 0 \leq i_j < d_j$ ，则计算过程用 `numpy` 语法可以写作：  
 $output[i_0, i_1, \dots, i_{p-1}, index[i_0, i_1, \dots, i_{p-1}, i_p, i_{p+1}, \dots, i_{r-1}], i_{p+1}, \dots, i_{r-1}] = data[i_0, i_1, \dots, i_{p-1}, i_p, i_{p+1}, \dots, i_{r-1}]$
- 对于上面的范例代码，就是：

```

index[0, 0, 0, 0] = 0 ⇒ output[0, 0, 0, 0] = update[0, 0, 0, 0] = -0.
index[0, 0, 0, 1] = 1 ⇒ output[0, 0, 1, 1] = update[0, 0, 0, 1] = -1.
index[0, 0, 0, 2] = 2 ⇒ output[0, 0, 2, 2] = update[0, 0, 0, 2] = -2.
...
index[0, 0, 1, 0] = 1 ⇒ output[0, 0, 1, 0] = update[0, 0, 1, 0] = -5.
index[0, 0, 1, 1] = 2 ⇒ output[0, 0, 2, 1] = update[0, 0, 2, 1] = -6.
index[0, 0, 1, 2] = 3 ⇒ output[0, 0, 3, 2] = update[0, 0, 3, 2] = -7.
...
index[0, 1, 0, 0] = 0 ⇒ output[0, 1, 0, 0] = update[0, 1, 0, 0] = -20.
index[0, 1, 0, 1] = 1 ⇒ output[0, 1, 1, 1] = update[0, 1, 0, 1] = -21.
index[0, 1, 0, 2] = 2 ⇒ output[0, 1, 2, 2] = update[0, 1, 0, 2] = -22.
...

```

- 计算公式恰好为 `GatherElement` 算子公式等号左右两项的索引进行交换
- `output` 元素的更新没有次序保证。如果两次更新指向 `output` 同一位置，且两次更新的值不同，则不能保证 `output` 该位置上的值选哪一次更新的结果。例如，将范例代码中 `data1` 改为 `data1 = np.zeros([nIn, cIn, hIn, wIn], dtype=np.int32)`，那么 `output[:, :, 0, :]` 值会是来自 `update` 的负的随机整数

## ND 模式

```

import numpy as np
from cuda import cudart
import tensorrt as trt

np.random.seed(97)
data0 = np.arange(2 * 3 * 4 * 5, dtype=np.float32).reshape(2, 3, 4, 5) # 输入数据
data1 = np.array([[[0, 2, 1, 1], [1, 0, 3, 2], [0, 1, 2, 3]], [[1, 2, 1, 1], [0, 0, 3, 2], [1, 1, 2, 3]]], dtype=np.int32)
data2 = -np.arange(2 * 3, dtype=np.float32).reshape(2, 3)

np.set_printoptions(precision=8, linewidth=200, suppress=True)
cudart.cudaDeviceSynchronize()

def scatterCPU(data0, data1, data2):
    output = data0
    for i in range(2):
        for j in range(3):
            print("i=%d, j=%d, index=%s, updateValue=%f"%(i, j, data1[i, j], data2[i, j]))
            output[data1[i, j][0], data1[i, j][1], data1[i, j][2], data1[i, j][3]] = data2[i, j]
    return output

logger = trt.Logger(trt.Logger.ERROR)
builder = trt.Builder(logger)
network = builder.create_network(1 << int(trt.NetworkDefinitionCreationFlag.EXPLICIT_BATCH))
config = builder.create_builder_config()
config.max_workspace_size = 1 << 30
inputT0 = network.add_input('inputT0', trt.DataType.FLOAT, (2, 3, 4, 5))
inputT1 = network.add_input('inputT1', trt.DataType.INT32, (2, 3, 4))
inputT2 = network.add_input('inputT2', trt.DataType.FLOAT, (2, 3))

```

```

#-----# 替换部分
scatterLayer = network.add_scatter(inputT0, inputT1, inputT2, trt.ScatterMode.ND)
#-----# 替换部分
network.mark_output(scatterLayer.get_output(0))
engineString = builder.build_serialized_network(network, config)
engine = trt.Runtime(logger).deserialize_cuda_engine(engineString)
context = engine.create_execution_context()
_, stream = cudart.cudaStreamCreate()

inputH0 = np.ascontiguousarray(data0.reshape(-1))
inputH1 = np.ascontiguousarray(data1.reshape(-1))
inputH2 = np.ascontiguousarray(data2.reshape(-1))
outputH0 = np.empty(context.get_binding_shape(3), dtype=trt.nptype(engine.get_binding_dtype(3)))
_, inputD0 = cudart.cudaMallocAsync(inputH0.nbytes, stream)
_, inputD1 = cudart.cudaMallocAsync(inputH1.nbytes, stream)
_, inputD2 = cudart.cudaMallocAsync(inputH2.nbytes, stream)
_, outputD0 = cudart.cudaMallocAsync(outputH0.nbytes, stream)

cudart.cudaMemcpyAsync(inputD0, inputH0.ctypes.data, inputH0.nbytes,
cudart.cudaMemcpyKind.cudaMemcpyHostToDevice, stream)
cudart.cudaMemcpyAsync(inputD1, inputH1.ctypes.data, inputH1.nbytes,
cudart.cudaMemcpyKind.cudaMemcpyHostToDevice, stream)
cudart.cudaMemcpyAsync(inputD2, inputH2.ctypes.data, inputH2.nbytes,
cudart.cudaMemcpyKind.cudaMemcpyHostToDevice, stream)
context.execute_async_v2([int(inputD0), int(inputD1), int(inputD2), int(outputD0)], stream)
cudart.cudaMemcpyAsync(outputH0.ctypes.data, outputD0, outputH0.nbytes,
cudart.cudaMemcpyKind.cudaMemcpyDeviceToHost, stream)
cudart.cudaStreamSynchronize(stream)

print("inputH0 :", data0.shape)
print(data0)
print("inputH1 :", data1.shape)
print(data1)
print("inputH2 :", data2.shape)
print(data2)
print("outputH0:", outputH0.shape)
print(outputH0)

cudart.cudaStreamDestroy(stream)
cudart.cudaFree(inputD0)
cudart.cudaFree(outputD0)

resCPU = scatterCPU(data0,data1,data2)
print("diff:\n",outputH0 - resCPU)

```

- 输入张量 0 形状 (2, 3, 4, 5)

$$\left[ \begin{array}{c} \left[ \begin{array}{ccccc} 0. & 1. & 2. & 3. & 4. \\ 5. & 6. & 7. & 8. & 9. \\ 10. & 11. & 12. & 13. & 14. \\ 15. & 16. & 17. & 18. & 19. \end{array} \right] \left[ \begin{array}{ccccc} 20. & 21. & 22. & 23. & 24. \\ 25. & 26. & 27. & 28. & 29. \\ 30. & 31. & 32. & 33. & 34. \\ 35. & 36. & 37. & 38. & 39. \end{array} \right] \left[ \begin{array}{ccccc} 40. & 41. & 42. & 43. & 44. \\ 45. & 46. & 47. & 48. & 49. \\ 50. & 51. & 52. & 53. & 54. \\ 55. & 56. & 57. & 58. & 59. \end{array} \right] \\ \left[ \begin{array}{ccccc} 60. & 61. & 62. & 63. & 64. \\ 65. & 66. & 67. & 68. & 69. \\ 70. & 71. & 72. & 73. & 74. \\ 75. & 76. & 77. & 78. & 79. \end{array} \right] \left[ \begin{array}{ccccc} 80. & 81. & 82. & 83. & 84. \\ 85. & 86. & 87. & 88. & 89. \\ 90. & 91. & 92. & 93. & 94. \\ 95. & 96. & 97. & 98. & 99. \end{array} \right] \left[ \begin{array}{ccccc} 100. & 101. & 102. & 103. & 104. \\ 105. & 106. & 107. & 108. & 109. \\ 110. & 111. & 112. & 113. & 114. \\ 115. & 116. & 117. & 118. & 119. \end{array} \right] \end{array} \right]$$

- 输入张量 1 形状 (2, 3, 4)

$$\left[ \begin{bmatrix} 0 & 2 & 1 & 1 \\ 1 & 0 & 3 & 2 \\ 0 & 1 & 2 & 3 \end{bmatrix} \begin{bmatrix} 1 & 2 & 1 & 1 \\ 0 & 0 & 3 & 2 \\ 1 & 1 & 2 & 3 \end{bmatrix} \right]$$

- 输入张量 2 形状 (2, 3)

$$\begin{bmatrix} -0. & -1. & -2. \\ -3. & -4. & -5. \end{bmatrix}$$

- 输出张量 0 形状 (2, 3, 4, 5)

$$\left[ \begin{bmatrix} \begin{bmatrix} 0. & 1. & 2. & 3. & 4. \\ 5. & 6. & 7. & 8. & 9. \\ 10. & 11. & 12. & 13. & 14. \\ 15. & 16. & -4. & 18. & 19. \end{bmatrix} \begin{bmatrix} 20. & 21. & 22. & 23. & 24. \\ 25. & 26. & 27. & 28. & 29. \\ 30. & 31. & 32. & -2. & 34. \\ 35. & 36. & 37. & 38. & 39. \end{bmatrix} \begin{bmatrix} 40. & 41. & 42. & 43. & 44. \\ 45. & -0. & 47. & 48. & 49. \\ 50. & 51. & 52. & 53. & 54. \\ 55. & 56. & 57. & 58. & 59. \end{bmatrix} \end{bmatrix} \right]$$

$$\left[ \begin{bmatrix} \begin{bmatrix} 60. & 61. & 62. & 63. & 64. \\ 65. & 66. & 67. & 68. & 69. \\ 70. & 71. & 72. & 73. & 74. \\ 75. & 76. & -1. & 78. & 79. \end{bmatrix} \begin{bmatrix} 80. & 81. & 82. & 83. & 84. \\ 85. & 86. & 87. & 88. & 89. \\ 90. & 91. & 92. & -5. & 94. \\ 95. & 96. & 97. & 98. & 99. \end{bmatrix} \begin{bmatrix} 100. & 101. & 102. & 103. & 104. \\ 105. & -3. & 107. & 108. & 109. \\ 110. & 111. & 112. & 113. & 114. \\ 115. & 116. & 117. & 118. & 119. \end{bmatrix} \end{bmatrix} \right]$$

- 范例代码 2，修改上面代码的若干行：

```
data0 = np.arange(2 * 3 * 4 * 5, dtype=np.float32).reshape(2, 3, 4, 5)
data1 = np.array([[0, 2, 1], [1, 0, 3], [0, 1, 2], [1, 2, 1], [0, 0, 3], [1, 1, 2]], dtype=np.int32)
data2 = -np.arange(6*5, dtype=np.float32).reshape(6,5)

def scatterCPU(data0,data1,data2):
    output = data0
    for i in range(6):
        print("i=%d,index=%s,updateValue=%s"%(i,data1[i]),data2[i])
        output[data1[i][0],data1[i][1],data1[i][2]] = data2[i]
    return output

inputT0 = network.add_input('inputT0', trt.DataType.FLOAT, (2, 3, 4, 5))
inputT1 = network.add_input('inputT1', trt.DataType.INT32, (6, 3))
inputT2 = network.add_input('inputT2', trt.DataType.FLOAT, (6, 5))
```

- 含义：参考 [Onnx ScatterND 算子](#) 和 [TensorRT C++ API 说明](#)
- 数据张量 data 形状  $[d_0, d_1, \dots, d_{r-1}]$  (r 维)，索引张量 index 形状  $[a_0, a_1, \dots, a_{q-1}]$  (q 维)，更新张量 update 形状  $[b_0, b_1, \dots, b_{s-1}]$  ( $s = r - a_{q-1} + q - 1$  维)，输出张量 output 形状与 data 相同
- 若  $r = a_{q-1}$  (范例代码 1)，则  $s = q - 1$ ，此时对于  $0 \leq i < q$ ，有  $b_i = a_i$  (update 比 index 少一维，且各维尺寸跟 index 去掉最低维后对应相等)
- 若  $r > a_{q-1}$  (范例代码 2)，则  $s > q - 1$ ，此时对于  $0 \leq i < q$ ，有  $b_i = a_i$ ，对于  $q \leq i < s$ ，有  $b_i = d_{a_{q-1}+i-q}$  (也即  $b_q = d_{a_{q-1}}, b_{q+1} = d_{a_{q-1}+1}, \dots$ ，最后一项  $i = s - 1$ ，此时  $b_{s-1} = d_{a_{q-1}+s-1-q} = d_{r-1}$ ，恰好取到 data 的最后一维的尺寸)
- 用 numpy 语法，记

```
q = len(index.shape)
nIndex = np.prod(index.shape[:-1])
index2D = index.reshape(nIndex, index.shape[-1])
update2D = update.reshape(nIndex, *update.shape[q-1:])
```

- 那么计算结果可以表示为

```
for i in nIndex:
    output[*index2D[i]] = update2D[i]
```

- 对于上面的范例代码 1，就是：

$$i = 0 \Rightarrow index2D[0] = [0, 2, 1, 1] \Rightarrow output[0, 2, 1, 1] = update2D[0] = -0.$$

...

$$i = 5 \Rightarrow index2D[5] = [1, 1, 2, 3] \Rightarrow output[1, 1, 2, 3] = update2D[5] = -5.$$

- 或者还原回 index 和 update 的原始下标来表示，就是：

$$i = 0, j = 0 \Rightarrow index[0, 0] = [0, 2, 1, 1] \Rightarrow output[0, 2, 1, 1] = update[0, 0] = -0.$$

$$i = 0, j = 1 \Rightarrow index[0, 1] = [1, 0, 3, 2] \Rightarrow output[1, 0, 3, 2] = update[0, 1] = -1.$$

$$i = 0, j = 2 \Rightarrow index[0, 2] = [0, 1, 2, 3] \Rightarrow output[0, 1, 2, 3] = update[0, 2] = -2.$$

...

$$i = 1, j = 0 \Rightarrow index[1, 0] = [1, 2, 1, 1] \Rightarrow output[1, 2, 1, 1] = update[1, 0] = -3.$$

$$i = 1, j = 1 \Rightarrow index[1, 1] = [0, 0, 3, 2] \Rightarrow output[0, 0, 3, 2] = update[1, 1] = -4.$$

$$i = 1, j = 2 \Rightarrow index[1, 2] = [1, 1, 2, 3] \Rightarrow output[1, 1, 2, 3] = update[1, 2] = -5.$$

- 对于上面的范例代码 2，就是：

$$i = 0 \Rightarrow index2D[0] = [0, 2, 1] \Rightarrow output[0, 2, 1] = update2D[0] = [-0., -1., -2., -3., -4.]$$

...

$$i = 5 \Rightarrow index2D[5] = [1, 1, 2] \Rightarrow output[1, 1, 2] = update2D[5] = [-25., -26., -27., -28., -29.,]$$

- 说明：

- 记  $nIndex = a_0 * a_1 * \dots * a_{q-2}$ ,
- 把 index 变形为  $nIndex$  行  $a_{q-1}$  列的矩阵 index2D，用其每一行来索引 data，同时把 update 变形为 nIndex 组形状为  $[b_{q-1}, \dots, b_{s-1}]$  的张量 update2D
- 如果  $r = a_{q-1}$ （范例代码 1），那么 index 的第  $i$  行作为索引恰好取到 output（或 data）的一个元素（ $np.shape(output[*index2D[i]]==[])$ ）；而此时  $b_{q-1} = b_{s-1} = 1$ （因为 update 只有  $q-1$  维，全在  $nIndex$  维度上了），该索引在 update2D 中也索引到一个元素，于是使用 update2D 的该元素来替换 output 的对应元素
- 如果  $r > a_{q-1}$ （范例代码 2），记  $nD = r - a_{q-1}$ ，那么 index 的第  $i$  行作为索引会取到 output（或 data）的一个 nD 维子张量（ $len(np.shape(output[*index2D[i]]))==nD$ ），形状  $[d_{a_{q-1}}, \dots, d_{r-1}]$ ；此时该索引在 update2D 中也索引到一个 nD 维的子张量，形状也是  $[d_{a_{q-1}}, \dots, d_{r-1}]$ ，于是使用 update2D 的该元素来替换 output 的对应元素
- 不满足  $s = r - a_{q-1} + q - 1$  时报错：

[TRT] [E] 4: [graphShapeAnalyzer.cpp::computeOutputExtents::1032] Error Code 4: Miscellaneous ((Unnamed Layer\* 0) [Scatter]: error while lowering shape of node)

- $b_i$  不满足约束条件时报错：

[TRT] [E] 4: [graphShapeAnalyzer.cpp::processCheck::581] Error Code 4: Internal Error ((Unnamed Layer\* 0) [Scatter]: dimensions not compatible for ScatterND)